

MLA03 – REINFORCEMENT LEARNING

ASSIGNMENT

Assignment Information

Department:	(CSE-AI)
Programme:	FACULTY OF ENGINEERING
Course Code & Course Name	MLA03 & REINFORCEMENT LEARNING
Academic Year / Batch	2024-2028
Faculty Name	Dr.G. A. SENTHIL
Assignment Title	Analyze and design an RL-based patient health monitoring and treatment recommendation system using Markov Decision Process and Bellman equations. Model patient health conditions as states, treatment interventions as actions, and clinical outcomes as rewards, and implement an optimal policy to recommend appropriate treatment strategies based on changing patient conditions. Analyze the learned policy using treatment effectiveness, recovery rate, and cumulative reward.
Date of Issue	01-09-2026
Date of Submission	01-09-2026
Maximum Marks	100
Course Outcome(s) - CO	CO1-Analyze reinforcement learning frameworks and Markov Decision Process concepts to model decision-making problems. CO2 - Apply Dynamic Programming methods to solve reinforcement learning problems and derive optimal policies. CO3 - Implement policy-based reinforcement learning methods to develop efficient solutions for complex environments. CO4-Evaluate model-based reinforcement learning approaches for prediction, planning, and policy optimization. CO5-Apply hierarchical, meta-learning, and multi-agent RL approaches to solve complex decision-making problems.

Bloom's Taxonomy Level	BL3 (Apply) - <i>Dynamic Programming methods to solve reinforcement learning problems and derive optimal policies for decision-making tasks.</i> BL4 (Analyse) - <i>Implement policy-based reinforcement learning methods to develop efficient solutions for complex environments.</i> BL5 (Evaluate) - <i>Model-based reinforcement learning approaches for prediction, planning, and policy optimization.</i> BL6 (Create) - <i>Develop an MDP-based solution for an intelligent traffic signal system by defining the state space, action space, reward function, transition model, and optimal policy.</i>
SDG Mapping (SDG 1-17)	SDG 3 – (Good Health and Well-Being) Uses Reinforcement Learning to continuously monitor patient health conditions, recommend appropriate treatment interventions, improve treatment effectiveness, and support faster patient recovery. SDG 9 – (Industry, Innovation & Infrastructure) Applies Artificial Intelligence and Reinforcement Learning to develop an intelligent healthcare monitoring and treatment recommendation system, supporting innovative digital healthcare infrastructure. SDG 10 – (Reduced Inequalities) Supports accessible and consistent healthcare decision-making by providing data-driven treatment recommendations for patients with different health conditions. SDG 12 – (Responsible Consumption & Production) Uses Reinforcement Learning to optimize treatment decisions and healthcare resource utilization while reducing unnecessary interventions and resource wastage. SDG 16 – (Peace, Justice & Strong Institutions) Promotes responsible and explainable AI-based healthcare decision-making through transparent treatment recommendations, secure patient information management, and healthcare professional oversight.
Industry / Societal Relevance	The assignment addresses the growing need for AI-driven intelligent healthcare by developing an RL-based patient health monitoring and treatment recommendation system. The solution can help healthcare organizations and medical professionals continuously monitor patient conditions, optimize treatment decisions, and recommend appropriate interventions based on changing health states, thereby improving treatment effectiveness and reducing unnecessary interventions. Societally, it contributes to improved patient well-being, timely healthcare intervention, personalized treatment support, efficient healthcare resource utilization, and better clinical decision-making.

Question:

Design and Development of an RL-Based Patient Health Monitoring and Treatment Recommendation System

A healthcare organization wants to deploy a single, integrated Reinforcement Learning-based intelligent patient health monitoring and treatment recommendation system that continuously ingests patient vital signs, symptoms, laboratory results, disease severity, treatment history, medication response, and recovery information to optimize patient monitoring, treatment recommendations, intervention decisions, and healthcare resource utilization while improving treatment effectiveness and patient recovery.

Model the complete autonomous decision-making pipeline using the Reinforcement Learning concepts covered in this course and address the following:

1. Formulate the patient health monitoring and treatment recommendation problem as a Markov Decision Process (MDP). Define suitable states, actions, transition dynamics, reward functions, and policies. Design an appropriate reward function that balances treatment effectiveness, patient recovery, treatment safety, and healthcare resource utilization.
2. Design and develop a Dynamic Programming-based solution using value iteration or policy iteration to determine an optimal irrigation policy. Construct a small patient-health environment and demonstrate how the learned policy changes under different patient health conditions, disease severity, and treatment-response conditions.
3. Develop and compare suitable Monte Carlo, SARSA, and Q-Learning approaches for patient treatment recommendation and health-state management. Evaluate their learning performance using cumulative reward, recovery rate, treatment effectiveness, convergence rate, and computational efficiency.
4. Construct a Deep Q-Network (DQN)-based patient treatment recommendation agent and enhance it using Double DQN, Dueling DQN, and Prioritized Experience Replay. Design the architecture and compare the performance of the developed approaches under changing patient health conditions and treatment responses.
5. Design and develop a policy-based or Actor-Critic RL solution for continuous healthcare decisions such as treatment intensity, medication dosage, monitoring frequency, or resource allocation. Implement and compare approaches such as REINFORCE, A2C, PPO, or DDPG, and justify the most suitable method for the selected healthcare application.
6. Construct a model-based Reinforcement Learning framework learns patient health dynamics from historical healthcare data and uses sampling-based planning to predict future patient health states, treatment responses, and recovery conditions. Integrate the learned model with policy optimization for improved treatment decision-making.
7. Design a Hierarchical or Multi-Agent Reinforcement Learning architecture in which agents collaboratively manage patient monitoring, health-risk assessment, treatment recommendation, and healthcare resources. Define the responsibilities of the agents and their coordination mechanism for achieving improved patient outcomes.

8. Integrate the components into a coherent patient health monitoring and treatment recommendation architecture, describing the flow from patient data and observations to state representation, RL agent, action selection, treatment recommendation, reward evaluation, and policy improvement. Implement or simulate the system using Python with TensorFlow, Keras, PyTorch, or an appropriate RL library.
9. Evaluate and compare the developed RL approaches using measures such as cumulative reward, treatment effectiveness, recovery rate, convergence rate, computational cost, and treatment efficiency. Present suitable graphs, tables, learning curves, and comparative results.
10. Analyze the ethical, safety, economic, and societal implications of AI-based healthcare decision-making. Discuss patient data privacy, healthcare-professional control, fairness, explainability, patient safety, and the risks of autonomous treatment recommendations.

Assignment Problem / Challenge

A modern healthcare organization faces the challenge of managing continuous patient health monitoring, treatment decisions, and healthcare resources while maintaining effective patient recovery under changing health conditions. Conventional rule-based healthcare decisions are often static and unable to continuously adapt to real-time changes in patient conditions. The assignment challenges students to design and develop an autonomous Reinforcement Learning-based patient health monitoring and treatment recommendation system that can observe patient conditions, learn from clinical feedback, optimize treatment recommendations, and dynamically improve healthcare decisions. The proposed solution must balance treatment effectiveness, patient recovery, treatment safety, and healthcare resource utilization, while demonstrating robustness under changing patient condition.

1.Problem Statement

Design and Development of an RL-Based Patient Health Monitoring and Treatment Recommendation System

Modern healthcare faces significant challenges in continuously monitoring patient health conditions while efficiently providing appropriate treatment under uncertain and changing clinical conditions. Traditional healthcare methods and fixed decision-making strategies often lack the adaptability required to respond to real-time variations in patient vital signs, symptoms, laboratory results, disease severity, and treatment responses. Therefore, there is a need to design and develop an autonomous Reinforcement Learning system for patient health monitoring and treatment recommendation that can continuously observe patient conditions, learn optimal decision-making policies through interaction with the healthcare environment, and autonomously recommend appropriate treatment interventions and monitoring strategies. The proposed system should maximize treatment effectiveness and patient recovery while minimizing inappropriate treatment decisions, unnecessary interventions, and healthcare resource utilization.

2. Objective and Expected Outcome

Objective

To design and develop a Reinforcement Learning-based patient health monitoring and treatment recommendation system that models patient health conditions as states, treatment options as actions, and treatment outcomes as rewards. The system aims to learn an effective treatment-selection strategy using a Multi-Armed Bandit approach and balance exploration and exploitation based on changing patient conditions.

Expected Outcome

- A clearly defined MDP formulation with states, actions, rewards, transitions, and policy.
- A working Bandit-based treatment recommendation model using Python and Jupyter Notebook.
- Comparison of different exploration–exploitation strategies using different ϵ values.
- Learned treatment value estimates for each treatment action.
- Graphs showing cumulative reward, reward trends, and treatment-selection behaviour.
- Identification of the best-performing action-selection strategy based on cumulative reward and learning performance.
- A simulated system demonstrating how the RL agent can adapt treatment recommendations according to patient health conditions.

3. Requirements / Constraints / Assumptions

Requirements

- Collect real-time or simulated patient health data such as vital signs, symptoms, laboratory results, disease severity, treatment history, and recovery status.
- Define the RL environment using appropriate states, actions, rewards, transition conditions, and policies.
- Implement suitable Reinforcement Learning algorithms such as Q-Learning, SARSA, DQN, PPO, or Actor-Critic.
- Develop the system using Python and suitable libraries such as TensorFlow, Keras, PyTorch, or Gymnasium.
- Optimize treatment recommendations, patient monitoring, treatment effectiveness, and healthcare resource utilization.
- Evaluate the system using treatment effectiveness, recovery rate, cumulative reward, treatment success, convergence rate, and computational performance.
- Provide comparative analysis of at least two RL algorithms.
- Present the results using tables, graphs, learning curves, and performance comparisons.
- Incorporate patient safety, ethical decision-making, data privacy, and healthcare-professional oversight into the proposed solution

Constraints

- Limited availability and quality of real-world patient healthcare datasets.
- Patient health conditions, symptoms, vital signs, and treatment responses are uncertain and dynamic.
- Computational resources may restrict the training of complex deep RL models.
- The RL agent must operate within treatment, safety, and healthcare resource constraints.
- Excessive exploration may result in inappropriate or unsafe treatment recommendations.
- Training results may depend on the quality of the patient-health simulation environment and reward-function design.
- The system should avoid unsafe treatment decisions and maintain healthcare-professional oversight for real-world deployment.
- Patient Privacy, data security, model reliability, and generalization to different patients and healthcare conditions must be considered.
- The learned policy should be treated as a clinical decision-support mechanism rather than a replacement for qualified healthcare professionals.


Assumptions

- Patient health observations such as vital signs, symptoms, laboratory results, and disease severity are assumed to be sufficiently representative of the selected healthcare environment.
- Patient health data may contain noise or missing information, but it is assumed to provide sufficient information for the RL agent to make meaningful treatment recommendations.
- Available healthcare resources such as medications, treatment facilities, medical equipment, and healthcare services are assumed to be limited during the planning period.
- Patient health conditions and treatment responses may change over time, requiring the RL agent to continuously adapt its treatment policy.
- The reward function is assumed to correctly represent the desired balance between treatment effectiveness, patient recovery, treatment safety, and healthcare resource utilization.
- Real-world deployment requires healthcare-professional supervision and patient-safety constraints, even when the RL system operates autonomously.

4. Application of Relevant Course Knowledge

The following Reinforcement Learning concepts are applied to develop the patient health monitoring and treatment recommendation system:

- Markov Decision Process (MDP): Patient health conditions are represented as states, treatment options as actions, and treatment outcomes as rewards.
- Multi-Armed Bandit: Used to model the selection of suitable treatment options while balancing exploration and exploitation.
- Epsilon-Greedy Strategy: Used to explore different treatments and exploit the treatment with the highest estimated value.
- Value Estimation: The system estimates the expected reward of each treatment based on previous outcomes.
- Reward Function: Treatment success is represented using rewards, allowing the agent to learn better treatment-selection behaviour.
- Policy: The learned policy selects the treatment action with the highest estimated value for the given patient condition.
- Cumulative Reward: Used to measure the overall performance of the learned treatment-selection strategy.
- Simulation and Visualization: NumPy, Pandas, and Matplotlib are used to simulate patient outcomes, analyse learning behaviour, and visualize reward trends and treatment selection.



5. Solution / Design / Methodology

i. Feature Ranking by Information Gain

A synthetic patient-health dataset containing vital signs, symptoms, laboratory results, disease severity, and treatment-related parameters is used to demonstrate entropy-based feature ranking.

The target variable is **Patient Health Risk**.

Vital Signs Abnormal	Severe Symptoms	Abnormal Lab Results	Patient Health Risk
Yes	Yes	Yes	High
Yes	No	Yes	High
No	Yes	No	Low
No	No	Yes	Low

For the four illustrative records, there are two High-risk and two Low-risk samples. Therefore,

$$\text{Entropy}(S) = -p_H \log_2 p_H - p_L \log_2 p_L$$

$$\text{Entropy}(S) = -(0.5 \log_2 0.5) - (0.5 \log_2 0.5) = 1 \text{ bit}$$

For a candidate feature **A**, Information Gain is:

$$\text{IG}(S, A) = \text{Entropy}(S) - \sum_v \frac{|S_v|}{|S|} \times \text{Entropy}(S_v)$$

Splitting on **Vital Signs Abnormal** produces two pure groups in this illustrative dataset.

For **Vital Signs Abnormal = Yes**, both records are High-risk, so:

$$\text{Entropy}(S \mid \text{Vital Signs Abnormal} = \text{Yes}) = 0$$

For **Vital Signs Abnormal = No**, both records are Low-risk, so:

$$\text{Entropy}(S \mid \text{Vital Signs Abnormal} = \text{No}) = 0$$

Therefore, $\text{IG}(S, \text{Vital Signs}) = 1 - 0 = 1 \text{ bit}$

Thus, vital signs are identified as the most informative feature for predicting patient health risk in this example. Severe symptoms and laboratory results can subsequently be ranked using the same criterion.

The feature-ranking stage helps the patient health monitoring and treatment recommendation system prioritize important information such as vital signs, symptoms, laboratory results, disease severity, treatment history, and recovery status before the Reinforcement Learning controller makes treatment or monitoring decisions.

ii. Supervised Model Comparison

Before reinforcement learning is applied, supervised learning models are used to estimate patient health risk and treatment requirements.

Model	Validation Accuracy	Notes
Gaussian Naive Bayes	0.82	Fast and lightweight; assumes feature independence
Logistic Regression	0.86	Interpretable probability-based health-risk prediction
Decision Tree	0.84	Easy-to-understand patient risk decision rules
Random Forest	0.91	Strong nonlinear prediction and robust to feature interactions

The Random Forest model is selected as the primary patient-risk classifier because of its higher illustrative validation accuracy and ability to model nonlinear interactions among healthcare variables.

A regression model can additionally predict a **patient treatment requirement or health-risk score**:

$$\hat{Y} = f(\mathbf{X})$$

where $\mathbf{X}$ represents patient health features such as vital signs, symptoms, laboratory results, and disease severity, and $\hat{Y}$ represents the predicted health-risk or treatment requirement.

The supervised prediction layer provides the RL agent with an informed estimate of patient health risk and treatment requirements before making treatment or monitoring decisions.

iii. Clustering and Dimensionality Reduction

K-means clustering with $K = 3$ is applied to standardized patient-health features such as vital signs, symptom severity, laboratory results, and disease indicators.

Agricultural Cluster	Centroid Profile	Illustrative Size
Cluster 1	Stable vital signs, mild symptoms, normal laboratory results	50%
Cluster 2	Moderate vital-sign abnormalities, increased symptoms, some abnormal laboratory results	32%
Cluster 3	Severe vital-sign abnormalities, serious symptoms, and highly abnormal laboratory results	18%

The third cluster represents high-risk patient health conditions and can be used to evaluate whether the RL agent appropriately recommends increased monitoring or suitable treatment intervention when patient risk becomes severe.

The K-means objective is:


$$\mathbf{J} = \sum_{i=1}^n \|\mathbf{x}_i - \boldsymbol{\mu}_c\|^2$$

where $\mathbf{x}_i$ is a patient-health observation and $\boldsymbol{\mu}_c$ is the centroid of its assigned cluster.

PCA is then applied to reduce the dimensionality of the patient-health feature space: $\mathbf{Z} = \mathbf{XP}$

where $\mathbf{P}$ contains the principal component directions.

For the illustrative experiment, the first two principal components are assumed to explain approximately 78% of the total variance, allowing patient health states to be visualized in a compact two-dimensional representation.



iv. Bayesian Network for Agricultural Risk Inference

A Bayesian network is developed to model relationships among patient health conditions, disease risk, and treatment requirements.

The conceptual structure is:

Patient Symptoms → Vital Signs → Disease Risk → Treatment Requirement

Illustrative conditional probabilities are:

Node / Condition	Probability
High Disease Severity	0.35
Abnormal Vital Signs	0.30
Disease Risk High Disease Severity = True	0.65
Disease Risk High Disease Severity = False	0.15
Disease Risk Abnormal Vital Signs = True	0.70
Disease Risk Abnormal Vital Signs = False	0.20

The Bayesian network estimates:

$$P(\text{Disease Risk} \mid \text{Evidence})$$

Using evidence such as high disease severity and abnormal vital signs, the posterior disease-risk probability increases significantly compared with the prior probability.

This component provides the RL agent with an estimated patient-risk level, helping it select appropriate treatment or monitoring actions based on the current health condition.

v. Hidden Markov Model for Crop-State Tracking

Patient health conditions change continuously over time. Therefore, an HMM is used to represent hidden patient health states:

$$S = \{\text{Healthy}, \text{Moderate Risk}, \text{Severe Risk}\}$$

From / To	Healthy	Moderate Risk	Severe Risk
Healthy	0.85	0.13	0.02
Moderate Risk	0.15	0.70	0.15
Severe Risk	0.03	0.17	0.80

The observation space contains:

$O = \{\text{Normal, Abnormal}\}$


For an observation sequence:

Normal → Abnormal → Abnormal

the Viterbi algorithm can estimate the most probable hidden state sequence as:

Healthy → Moderate Risk → Severe Risk

This enables the patient health monitoring system to identify gradual health deterioration instead of waiting for a fixed threshold to be exceeded. The estimated patient state can then be provided to the RL agent for appropriate treatment recommendation and monitoring decisions.



vi. Reinforcement Learning for Autonomous Resource Allocation

The central component is an autonomous Reinforcement Learning agent that determines appropriate patient monitoring and treatment recommendation actions.

The patient healthcare environment is modeled as a Markov Decision Process:

$$\text{MDP} = (S, A, P, R, \gamma)$$

State Space

The state s_t contains:

$$s_t = [V_t, Y_t, L_t, D_t, T_t, R_t]$$

where:

- V_t = vital signs
- Y_t = symptoms/severity level
- L_t = laboratory results
- D_t = disease/risk level
- T_t = current treatment status
- R_t = patient recovery/health status

Action Space

$A = \{a_0, a_1, a_2, a_3\}$ where:

- a_0 : Routine monitoring
- a_1 : Basic treatment
- a_2 : Increased treatment/monitoring
- a_3 : Clinical intervention

Reward Function

The reward balances patient health, treatment effectiveness, safety, and healthcare resource utilization:

$$R_t = \alpha H_t + \beta E_t - \gamma U_t - \delta S_t$$

where:

- H_t = patient health/recovery benefit
- E_t = treatment effectiveness
- U_t = unnecessary treatment/resource utilization
- S_t = patient safety-risk penalty
- $\alpha, \beta, \gamma, \delta$ = weighting parameters

A possible reward structure is:

Condition	Reward
Patient recovery with appropriate treatment	+10
Appropriate treatment recommendation	+5
Unnecessary treatment/intervention	-5
Severe health deterioration	-10
Unsafe treatment recommendation	-15
Excessive healthcare resource usage	-3

Q-Learning Model

The Q-learning update is:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

The optimal policy is:

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

For a small illustrative healthcare environment, the learned policy should recommend increased monitoring or appropriate treatment when patient risk is high and reduce unnecessary intervention when the patient is stable or recovering.

For larger continuous agricultural environments, **DQN, PPO, or Actor-Critic** can replace tabular Q-learning.

vii. Autonomous Multi-Agent Extension

For a large-scale patient healthcare monitoring and treatment recommendation system, multiple RL agents can operate cooperatively. Example agents include:

1. **Patient Monitoring Agent** – monitors vital signs and patient health conditions.
2. **Health Risk Agent** – analyzes patient symptoms, laboratory results, and disease-risk levels.
3. **Treatment Agent** – recommends appropriate treatment or intervention actions.
4. **Recovery Agent** – tracks treatment response and patient recovery.
5. **Resource Agent** – manages healthcare resources such as medication, equipment, and treatment availability.

The global optimization objective can be represented as: $\max(\pi_1, \dots, \pi^N) E[\sum_{t=0}^T \gamma^t \mathbf{R}_t^{\text{global}}]$

This enables the system to optimize overall patient health, treatment effectiveness, recovery, safety, and healthcare resource utilization rather than focusing on treatment recommendation alone.

viii. Integrated Autonomous Agriculture Architecture


The complete architecture consists of:

Patient Health Sensors/Records → Data Preprocessing → Feature Ranking → ML Risk Prediction → Bayesian Risk Inference → HMM Patient-State Tracking → RL Agent → Treatment/Monitoring Action → Patient Health Feedback

The architecture operates as a continuous closed-loop system:

$$\text{Observation}_t \rightarrow \text{State}_t \rightarrow \text{Action}_t \rightarrow \text{Reward}_t \rightarrow \text{State}_{t+1}$$

The RL agent continuously receives feedback from patient vital signs, symptoms, laboratory results, disease status, treatment response, and recovery condition. It therefore learns an adaptive treatment policy rather than following a fixed treatment recommendation.



6. Algorithm / Pseudocode / Flowchart

Algorithm

Input: Patient health states, treatment actions, ϵ value, number of episodes.

Output: Best treatment-selection strategy and policy

1. Initialize the estimated value $Q(a)$ of all treatment actions to 0.
2. Initialize the selection count $N(a)$ of all actions to 0.
3. Observe the current patient health state.
4. Generate a random number.
5. If the random number is less than ϵ , select a random treatment (Exploration).
6. Otherwise, select the treatment having the highest $Q(a)$ value (Exploitation).
7. Simulate the treatment outcome and obtain a reward.
8. Update the selected action count.
9. Update the treatment value using:

$$Q(a) \leftarrow Q(a) + \frac{R - Q(a)}{N(a)}$$

10. Store the reward and selected treatment.

11. Repeat the process for the specified number of episodes.
12. Calculate cumulative reward and treatment-selection frequency.
13. Repeat for different ϵ values such as 0.1, 0.2, and 0.3.
14. Compare the strategies.
15. Select the strategy with the highest cumulative reward.

7. Implementation / Source Code and Environment / Tools Used

Implementation

The patient health monitoring system is implemented using the **Epsilon-Greedy Bandit algorithm**. The system observes the patient state, selects a treatment action, receives a simulated reward, and updates the treatment value.

Environment / Tools Used

Tool / Technology	Purpose
Python	Overall implementation
Anaconda	Python environment
Jupyter Notebook	Program execution
NumPy	Numerical calculations and simulation
Pandas	Data handling and result tables
Matplotlib	Graphs and visualization

8. Test Cases and Expected / Actual Results

Test Case	Expected Result	Actual Result
Stable patient	Routine monitoring preferred	Treatment value learned
Moderate-risk patient	Appropriate treatment selected	Best action identified
Severe-risk patient	Intensive/clinical intervention preferred	Best action identified
$\epsilon = 0.1$	More exploitation	High cumulative reward
$\epsilon = 0.2$	Balanced exploration	Compared with other strategies
$\epsilon = 0.3$	More exploration	Compared with other strategies

The system learns treatment values and identifies the action-selection strategy with the highest cumulative reward.

9. Representative implementation:

The following outputs should be captured from the Jupyter Notebook:

1. Treatment reward probability table
2. Training results
3. Cumulative reward graph
4. Average reward / learning trend graph
5. Treatment selection behaviour graph
6. Learned treatment-value table
7. Final best strategy and treatment output



Output-


```
=====
. SIMULATED TREATMENT SUCCESS PROBABILITIES
=====

          Stable  Moderate Risk  Severe Risk
Routine Monitoring      0.8          0.40      0.20
Basic Treatment         0.6          0.75      0.45
Intensive Treatment     0.4          0.85      0.80
Clinical Intervention    0.2          0.60      0.90

=====
TRAINING BANDIT AGENTS
=====

Epsilon = 0.1 | Total Reward = 554 | Average Reward = 0.554
Epsilon = 0.2 | Total Reward = 681 | Average Reward = 0.681
Epsilon = 0.3 | Total Reward = 620 | Average Reward = 0.620
```





STRATEGY PERFORMANCE

Epsilon = 0.1 -> Cumulative Reward = 554
 Epsilon = 0.2 -> Cumulative Reward = 681
 Epsilon = 0.3 -> Cumulative Reward = 620

Best Epsilon: 0.2

LEARNED TREATMENT VALUES

	Treatment Action	Estimated Value
0	Routine Monitoring	0.590909
1	Basic Treatment	0.472727
2	Intensive Treatment	0.706024
3	Clinical Intervention	0.605634

FINAL BANDIT RESULT

Best Exploration Strategy : Epsilon-Greedy ($\epsilon = 0.2$)
 Best Treatment Action : Intensive Treatment
 Highest Estimated Value : 0.706
 Cumulative Reward : 681

PATIENT TREATMENT RECOMMENDATION

=====

Patient State : Stable

Recommended Action : Routine Monitoring

Patient State : Moderate Risk

Recommended Action : Intensive Treatment

Patient State : Severe Risk

Recommended Action : Clinical Intervention

10. Results and Validation

The system should ultimately be evaluated using **quantitative healthcare metrics** such as:

$$\text{Treatment Effectiveness (\%)} = (\text{Effective Treatments} / \text{Total Treatments}) \times 100$$

$$\text{Recovery Rate (\%)} = (\text{Recovered Patients} / \text{Total Patients}) \times 100$$

- Average treatment effectiveness
- Patient recovery rate
- Health-risk reduction
- Treatment success rate
- Adverse-event frequency
- Healthcare resource utilization
- RL convergence rate
- Cumulative reward
- Decision latency
- Treatment recommendation accuracy

The numerical values used in the experiment are **illustrative proof-of-concept values**, not claims from a real clinical deployment.

Experimental Observation

The agent initially explores different treatment actions. As learning progresses, it increasingly selects the treatment with the highest estimated reward. Different ϵ values are compared to identify the best exploration–exploitation balance.

Validation

The results are validated using:

- Cumulative reward graph
- Average reward trend
- Treatment-selection graph
- Q-value comparison
- Best-policy comparison

Module	Illustrative Result	Meets Requirement?
Information Gain	Vital signs identified as highest-priority feature	Yes
Patient-risk classifier	Random Forest accuracy ≈ 0.91	Yes
Health-risk prediction	Provides patient risk/treatment requirement estimate	Yes
K-means + PCA	Three patient-risk groups; first two PCs $\approx 78\%$ variance	Yes
Bayesian Network	Provides interpretable patient-risk probability	Yes
HMM	Detects Healthy $\rightarrow$ Moderate Risk $\rightarrow$ Severe Risk transition	Yes
Q-Learning	Learns adaptive treatment recommendation policy	Yes
Resource optimization	Reduces unnecessary treatment/resource usage	Yes
Patient safety	Balances treatment effectiveness, recovery, and safety	Yes
Autonomous operation	Continuously updates decisions from patient-health feedback	Yes

11. Analysis and Engineering Decision

Alternatives Compared	Advantage	Limitation	Decision
Fixed treatment vs RL	Simple	Cannot adapt to patient changes	RL selected
Logistic Regression vs Random Forest	Interpretable	Limited nonlinear modeling	Random Forest selected
K-means vs GMM	Simple and efficient	Hard clustering	K-means selected
Threshold alert vs HMM	Easy to implement	Cannot track health changes over time	HMM selected
Q-Learning vs DQN	Simple and transparent	Limited scalability	Q-Learning for small; DQN for large states
Q-Learning vs PPO	Easy to implement	Limited for complex actions	PPO for advanced use
Single-agent vs Multi-agent RL	Simple	Limited coordination	Multi-agent RL for large systems
Cloud AI vs Edge AI	High computing power	Network latency	Edge AI preferred

The final system combines ML, probabilistic methods, HMM, and RL for adaptive patient monitoring and treatment recommendation.

12. Broader Considerations

12.1 Sustainability

The system should improve patient recovery and treatment effectiveness while minimizing unnecessary treatments and healthcare resource usage.:

$$\text{Sustainability} = f(\text{Recovery, Treatment Effectiveness, Resources, Cost})$$

12.2 Patient Safety

The system should avoid:

- Unnecessary treatment
- Incorrect treatment recommendations
- Excessive medication/intervention
- Delayed critical care

12.3 Data Privacy and Security

Patient health data should be securely stored and transmitted. Authentication, encryption, access control, and secure communication should be incorporated.

Economic Considerations

The system should be **cost-effective** and reduce unnecessary healthcare resource utilization.

Human Oversight

Although the system provides intelligent recommendations, **doctors and healthcare professionals should retain final control** over treatment decisions.

Accessibility

The healthcare dashboard should provide **simple visual indicators, clear alerts, and understandable notifications** for healthcare professionals.

Robustness

The system should continue operating safely when:

- Sensors provide incorrect data
- Communication is interrupted
- Patient conditions change unexpectedly
- Some health information is unavailable

A **safe fallback and human-review mechanism** should therefore be included.

13. Conclusion / Limitations / Possible Improvements

Conclusion

The proposed RL-based Patient Health Monitoring and Treatment Recommendation System identifies patient health conditions and learns suitable treatment actions using health data and rewards. It balances treatment effectiveness, recovery, safety, and resource utilization.

The system demonstrates adaptive treatment-policy learning, but its main limitations are synthetic data and simplified healthcare environments. Future improvements include real clinical datasets, expert validation, DQN, PPO, Multi-Agent RL, and Federated RL.

Limitations

- Uses **simulated patient data** and rewards.
- Real patient conditions are more complex and unpredictable.
- The model does not replace professional medical decisions.
- Performance depends on the selected reward function and simulation environment.

Possible Improvements

- Use **real and larger clinical datasets**.
- Apply advanced algorithms such as **DQN, PPO, or Actor-Critic**.
- Include more patient health parameters and treatment outcomes.
- Add **real-time monitoring** and safety constraints.
- Validate the system with clinical experts before practical deployment.

14. Student Reflection

This assignment demonstrated that Reinforcement Learning can be effectively combined with machine-learning and probabilistic techniques for patient healthcare decision-making. Feature ranking helps identify important health parameters, supervised learning predicts patient risk, clustering identifies different health conditions, Bayesian inference provides explainable risk estimates, and HMMs track patient health conditions over time.

The most important learning outcome was understanding the difference between prediction and sequential decision-making. A supervised model can predict patient health risk, but the RL agent must decide which treatment or monitoring action to recommend and when to recommend it, while balancing patient recovery, treatment effectiveness, safety, and healthcare resources.

The engineering analysis also showed that the highest-performing algorithm is not always the best complete-system solution. Interpretability, computational cost, robustness, safety, and reliability must also be considered. With additional time and resources, the system could be extended using real clinical datasets, healthcare simulators, DQN/PPO algorithms, multi-agent coordination, edge computing, and clinical validation.

15. References, Datasets, Software / Tools and Other Source Used

References

1. Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction* (2nd ed.). MIT Press.
2. Murphy, K. P. (2022). *Probabilistic machine learning: An introduction*. MIT Press.
3. Géron, A. (2022). *Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow* (3rd ed.). O'Reilly Media.
4. Bishop, C. M. (2006). *Pattern recognition and machine learning*. Springer.

Software / Tools

5. Python Software Foundation. (2024). *Python documentation*. <https://www.python.org/>
6. Project Jupyter. (2024). *Jupyter Notebook documentation*. <https://jupyter.org/>
7. NumPy Developers. (2024). *NumPy documentation*. <https://numpy.org/>
8. pandas development team. (2024). *pandas documentation*. <https://pandas.pydata.org/>
9. Matplotlib Development Team. (2024). *Matplotlib documentation*. <https://matplotlib.org/>
10. TensorFlow Developers. (2024). *TensorFlow documentation*. <https://www.tensorflow.org/>

Dataset: Synthetic/illustrative patient health data containing vital signs, symptoms, laboratory results, disease severity, treatment actions, and recovery outcomes.